

PROFESSIONAL SUMMARY

Trilingual Research Engineer specializing in RL infrastructure and agentic evaluation at scale (rollout orchestration, multi-provider sandboxing, and grading harnesses) sustaining thousands of concurrent coding agents. Currently core R&D at *micro1* while pursuing an MS in Quantum Computing, with prior production ML experience in multilingual NLP, classification, and model evaluation across international markets.

TECHNICAL SKILLS

Languages : Python, TypeScript/JavaScript, Go, R, SQL, Java, Bash, HTML, CSS, $\LaTeX$
Frameworks : TensorFlow, PyTorch, Scikit-Learn/CuML, Apache Frameworks (Spark, Kafka, HBase, Hive, Hadoop, Phoenix, Airflow, Flume, Sqoop, NiFi, etc.), LGBM/XGBoost, SQLite3, SciPy-NumPy/CuPy, Pandas/CuDF, Matplotlib, Seaborn, Shapash
Cloud/Database Platforms : Jupyter Notebooks, MySQL/PostgreSQL, MongoDB, Microsoft Office Suite
DevOps and API Tools : Git/GitHub, Linux, Docker, TFX/TorchServe, Microsoft Azure, AWS, Google Cloud,
RL Infrastructure : E2B, Modal, Daytona, Harbor, rollout orchestration, SxS evaluation/harnessing

EXPERIENCE

- Member of Technical Staff, Research Engineering** | *micro1, Inc.*

December 2025 – Today

 - Core research engineer owning in-house RL rollout infrastructure on GCP/AWS, orchestrating 3.5M+ agent rollouts across E2B/Modal/Daytona peaking at 60k+ runs/day across multi-model agentic coding/non-coding pipelines.
 - Rebuilt the rollout stack twice in 3 months to keep pace with shifting customer and research requirements; migrated execution environments to daytona/Modal, cutting sandbox cold-start p50 111s → 23s (-79%) at sustained load.
 - Led R&D for multi-tenant RAG and internal data-auditing platforms and drove client-facing evals for Fortune 50 accounts.
- Member of Technical Staff, ML Engineering** | *Deledao Technologies Corp.*

August 2024 – December 2025

 - Spearheaded multilingual ML classification system/cross-department automation R&D for content classification and student wellness products, deploying multi-tiered models achieving 96%+ F1 score/98%+ accuracy/sub-400ms latency for Taiwanese market, built 10+ internal automations scaling infrastructure for 3x sales growth.
- Forward Deployed Engineer** | *MG Solutions*

March 2025 – Today

 - Migrated SMB clients from vendor-locked platforms (GoDaddy, Microsoft 365) to self-hosted VPS infrastructure over 2-5 month engagements, architecting Docker-based environments with Nginx reverse proxy, Cloudflare DDoS protection, and Prometheus monitoring; eliminated recurring SaaS costs while maintaining 99.5% uptime.
- Adjunct Instructor/Teaching Assistant** | *FTC-NC*

January 2024 – December 2025

 - Instructor for Linear Algebra. Collaborated with faculty to tutor Data Mining/Data Structures & Algorithm students, delivered coursework feedback and exam prep, and created QMD/ $\LaTeX$ lecture materials.

EDUCATION

Fei Tian College - Northern Campus

MS in Quantum Computing | GPA : N/A

September 2025 – December 2027

Fei Tian College - Northern Campus

BS in Data Science, magna cum laude | GPA : 3.80

January 2023 – May 2025

PROJECTS

- Dual-Layer Bilingual (EN/ZH) Web-Filter** | *Python – TensorFlow, CuML, CuPy, Jieba, Docker*

Jan – Nov 2025

 - Built and deployed a two-stage EN/ZH web filter for 13 sensitive categories in K-12 markets — token-frequency screening into a context-aware mmBERT transformer with per-category sensitivity tuning — accelerating training **30×** via RAPIDS/TensorFlow and containerizing an end-to-end TFX serving pipeline.
- Locality Sensitive Hashing (LSH) Recommendation System** | *Python - JAX*

May 2025 – Today

 - Authored an open-source, pip-installable LSH library (hyperplane hashing, JAX-parallelized scoring, Redis bucket store, Ruff/UV CI/CD) hitting ~15ms query latency at 1M scale, projected 228ms at 1B, deployed for RAG retrieval on resource-constrained hardware.

AWARDS AND CERTIFICATIONS

- Certificate of Department Superstar, AI Team** | *micro1*

May 2026
- Microsoft AI & ML Engineering** | *Microsoft*

February 2026
- AWS Cloud Solutions Architect** | *Amazon Web Services (AWS)*

January 2026
- IBM Deep Learning with PyTorch, Keras and Tensorflow** | *IBM*

November 2025
- (MLH) Best use of Google Gemini API** | *HackRPI 2025*

November 2025
- Advanced Machine Learning on Google Cloud** | *Google Cloud*

November 2024
- Natural Language Processing (NLP) Specialization** | *Deeplearning.AI*

November 2024
- Deep Learning Specialization** | *DeepLearning.AI*

October 2024
- IBM Data Engineering Professional Certificate** | *IBM*

September 2024
- 5x Dean’s List, 3x Department Academic Excellence Award** | *FTC-NC*

as of May 2025